

A stylized, light brown illustration of a plant with several leaves and a cluster of small, round fruits or berries, positioned on the left side of the slide.

DATABASE MANAGEMENT SYSTEM

Subash Manandhar

Chapter 7 : Query Processing

- Query processing is transforming a query written in high level language typically in SQL into correct and efficient execution strategy expressed in a low level language and to execute the strategy to retrieve required data.
- Basic steps in query processing:
 - Parsing and translation
 - Optimization
 - Evaluation




FIG: Steps in query processing

Chapter 7 : Query Processing

- **Parsing and Translating:**

- Translate the query into internal form. This is then translated into relational algebra.
- Parser checks syntax, verifies relation.
- The relational algebra representation of a query specifies only partially how to evaluate a query.
- A relational algebra expression may have many equivalent expressions.

- Consider a query:

SELECT salary FROM employee WHERE salary > 20000

Now, equivalent relational algebra expressions are

1. $\sigma_{\text{salary} > 20000}(\pi_{\text{salary}}(\text{employee}))$
2. $\pi_{\text{salary}}(\sigma_{\text{salary} > 20000}(\text{employee}))$

Chapter 7 : Query Processing

- **Evaluation:**

- The query execution engine takes a query evaluation plan, executes that plan and returns the answers to query.

- **Query Optimization:**

- Amongst all equivalent evaluation plans choose the one with lowest cost.
- Cost is estimated using statistical information from database catalog.

e.g. no. of tuples in each relation, size of tuples etc.


A sequence of primitive operations that can be used to evaluate a query is called query execution plan or query evaluation plan.

Once the query plan is chosen, the query is evaluated with that plan and result of query is output.

Chapter 7 : Query Processing

- **Construction of parse tree:**

- A leaf node is created for each base relation in query.
- A non leaf node is created for each intermediate relation produced by relational algebra expression.
- The root node represents result of query.
- The sequence of operation directed from the leaves to the root.
- E.g. $\pi_{\text{name, street}}(\sigma_{\text{status} = \text{'open'}}(\text{order} \bowtie \text{customer}))$



Chapter 7 : Query Processing

- **Query Optimization:**

- Is an important aspect of query processing.
- Is the activity of choosing an efficient execution strategy for processing a query.
- Aim is to choose one of many equivalent transformation that minimize resource usage.
- Reduce the total execution time of the query, which is the sum of execution times of all individual operations that make up query.
- There can be enormous differences in term of performance between different evaluation plans for same query.
- E.g. seconds Vs days to execute same query.

Chapter 7 : Query Processing

- **Cost based query Optimization:**

- Generate logically equivalent expressions by using a set of equivalence rules.
- Annotate the expressions to get alternative query evaluation plans.
- Select the cheapest plan based on estimated cost.
- Estimation of query evaluation cost based on statistical information from the catalog manager in combination with expected performance algorithms.

- **Importance of query Optimization:**

- Provides faster query processing
- Requires less cost per query
- Provides high performance of the system
- Consumes less memory

Chapter 7 : Query Processing

- **Measure of query cost:**

- Cost is generally measured as total elapsed time for answering query

- Many factors contribute to time cost
 - *disk accesses, CPU, or even network communication*
- Typically disk access is the predominant cost, and is also relatively easy to estimate. Measured by taking into account
 - Number of seeks * average-seek-cost
 - Number of blocks read * average-block-read-cost
 - Number of blocks written * average-block-write-cost
 - Cost to write a block is greater than cost to read a block
 - data is read back after being written to ensure that the write was successful
- For simplicity we just use the **number of block transfers** from disk and the **number of seeks** as the cost measures
 - t_T – time to transfer one block
 - t_S – time for one seek
 - Cost for b block transfers plus S seeks
$$b * t_T + S * t_S$$

Chapter 7 : Query Processing

- **Measure of query cost:**
 - We ignore CPU costs for simplicity
 - Real systems do take CPU cost into account
 - We do not include cost to writing output to disk in our cost formulae
- **Query operation:**
 - Selection Operation
 - The lowest level query processing operator for accessing data is the file scan.
 - search and retrieve records for a given selection condition.
 - Linear Search
 - Scan each file block and test all records to see whether they satisfy the selection condition.
 - Cost estimate = b_r block transfers + 1 seek
 - b_r denotes number of blocks containing records from relation r
 - If selection is on a key attribute, can stop on finding record
 - cost = $(b_r/2)$ block transfers + 1 seek
 - Linear search can be applied regardless of
 - selection condition or
 - ordering of records in the file, or
 - availability of indices

Chapter 7 : Query Processing

- **Selection Using Indices:**

- **Index scan** – search algorithms that use an index
 - selection condition must be on search-key of index.
- **A2 (primary index, equality on key).** Retrieve a single record that satisfies the corresponding equality condition
 - $\text{Cost} = (h_i + 1) * (t_T + t_S)$
- **A3 (primary index, equality on nonkey)** Retrieve multiple records.
 - Records will be on consecutive blocks
 - Let b = number of blocks containing matching records
 - $\text{Cost} = h_i * (t_T + t_S) + t_S + t_T * b$
- **A4 (secondary index, equality on nonkey).**
 - Retrieve a single record if the search-key is a candidate key
 - $\text{Cost} = (h_i + 1) * (t_T + t_S)$
 - Retrieve multiple records if search-key is not a candidate key
 - each of n matching records may be on a different block
 - $\text{Cost} = (h_i + n) * (t_T + t_S)$
 - Can be very expensive!

Chapter 7 : Query Processing

- **Sorting:**

- A query may specify that the output should be sorted.
- The processing of some relational query operations can be implemented more efficiently based on sorted relations (join operation)
- For relations that fit into memory, techniques like quick sort can be used.
- For relations that do not fit into memory an external merge sort can be used.


- **External Merge Sort**

- Let M denote memory size (in pages).
 - **Create sorted runs.**
 - Let i be 0 initially.
Repeatedly do the following till the end of the relation:
 - (a) Read M blocks of relation into memory
 - (b) Sort the in-memory blocks
 - (c) Write sorted data to run R_i ; increment i .
- Let the final value of i be N

Chapter 7 : Query Processing

- **External Merge Sort**
 - **Merge the runs (N-way merge).** We assume (for now) that $N < M$.
 - Use N blocks of memory to buffer input runs, and 1 block to buffer output. Read the first block of each run into its buffer page
 - **repeat**
 - Select the first record (in sort order) among all buffer pages
 - Write the record to the output buffer. If the output buffer is full write it to disk.
 - Delete the record from its input buffer page.
If the buffer page becomes empty **then**
read the next block (if any) of the run into the buffer.
 - **until** all input buffer pages are empty:
 - If $N \geq M$, several merge passes are required.
 - In each pass, contiguous groups of $M - 1$ runs are merged.
 - A pass reduces the number of runs by a factor of $M - 1$, and creates runs longer by the same factor.
 - E.g. If $M=11$, and there are 90 runs, one pass reduces the number of runs to 9, each 10 times the size of the initial runs
 - Repeated passes are performed till all runs have been merged into one.

Chapter 7 : Query Processing



Chapter 7 : Query Processing

- Equivalence Rule:

- Conjunctive selection operations can be deconstructed into a sequence of individual selections
 - $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$
- Selection operations are *commutative*
 - $\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$
- Cascade of projection operations (only final one)
 - $\pi_{A_1}(\pi_{A_2}(\dots(\pi_{A_n}(E))\dots)) = \pi_{A_1}(E)$
- Selections can be combined with cartesian products and theta joins
 - $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$
 - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$
- Theta join (and natural join) operations are *commutative*
 - $E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$
 - note that the order of attributes is ignored
- Natural join operations are *associative*
 - $(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$
- Theta joins are *associative* in the following manner
 - $(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$
 - where θ_2 contains attributes only from E_2 and E_3
- Union and intersection operations are *commutative*
 - $E_1 \cup E_2 = E_2 \cup E_1$
 - $E_1 \cap E_2 = E_2 \cap E_1$

Chapter 7 : Query Processing

- Equivalence Rule:

Union and intersection operations are *associative*

- $(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$
- $(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$

The *selection* operation *distributes* over *union*, *intersection* and *set difference*

- $\sigma_P(E_1 - E_2) = \sigma_P(E_1) - \sigma_P(E_2)$

The *projection distributes* over the *union* operation

- $\pi_A(E_1 \cup E_2) = (\pi_A(E_1)) \cup (\pi_A(E_2))$

Note that this is only a selection of equivalence rules

Chapter 7 : Query Processing

- **Heuristic Optimization:**

- Cost-based optimization is expensive, even with dynamic programming.
- Systems may use *heuristics* to reduce the number of choices that must be made in a cost-based fashion.
- Heuristic optimization transforms the query-tree by using a set of rules that typically (but not in all cases) improve execution performance:
 - Perform selection early (reduces the number of tuples)
 - Perform projection early (reduces the number of attributes)
 - Perform most restrictive selection and join operations before other similar operations.
 - Some systems use only heuristics, others combine heuristics with partial cost-based optimization.

Chapter 7 : Query Processing

- **Steps in Heuristic Optimization:**

- Deconstruct conjunctive selections into a sequence of single selection operations .
- Move selection operations down the query tree for the earliest possible execution .
- Execute first those selection and join operations that will produce the smallest relations .
- Replace Cartesian product operations that are followed by a selection condition by join operations .
- Deconstruct and move as far down the tree as possible lists of projection attributes, creating new projections where needed .
- Identify those sub trees whose operations can be pipelined, and execute them using pipelining.

Chapter 7 : Query Processing

- **Performance Tuning**

- **Database Statistics**

- Get the correct and updated statistics

- **Create Optimized Indexes**

- Have a right balance of indexes on tables

- **Specify optimizer hints in SELECT**

- specify the index name in SELECT query

- **Predetermine expected growth**

- specify an appropriate value for fill factor when creating indexes

- **Select limited data**

- Rather than filtering on the client, push as much filtering as possible on the server-end
 - Eliminate any obvious or computed columns

- **Drop indexes before loading data**

- drop the indexes on a table before loading a large batch of data. This makes the insert statement run faster.
 - Recreate the indexes after insertion.

- **Avoid foreign key constraints**

- Foreign keys constraints ensure data integrity at the cost of performance